

# Behavior Scoring - Project Documentation

Source: docs/api\_reference.md

## API Reference

### Overview

This reference describes the HTTP API implemented in app/main.py for the Behavior Scoring prototype, version 0.2.0. It is intended for local R&D use with synthetic, sample, or explicitly consented text only. It is not a production hiring-decision API.

Default base URL: `http://127.0.0.1:8000`

The currently mounted routes do not enforce authentication or authorization. POST `/api/auth/token` issues a development bearer token, but no route yet validates that token. Requests and JSON responses use `application/json` unless an endpoint explicitly returns a download.

### Endpoint Index

Method | Path | Purpose

GET | `/api/health` | Check backend and local Ollama availability.  
GET | `/api/rubric` | Return the active scoring rubric and exclusions.  
POST | `/api/score` | Score and persist one candidate profile.  
POST | `/api/webhooks` | Register an event webhook.  
POST | `/api/auth/token` | Exchange client credentials for a development bearer token.  
POST | `/api/candidates/import` | Persist up to 100 imported candidate profiles.  
GET | `/api/scores` | List persisted score runs with filters and pagination.  
GET | `/api/scores/analytics` | Return aggregate score and review metrics.  
GET | `/api/scores/export` | Download filtered score runs as CSV or JSON.  
GET | `/api/scores/compare` | Compare selected persisted score runs.  
GET | `/api/scores/{score_id}` | Return one persisted score run.  
PATCH | `/api/scores/{score_id}/review` | Record a human-review outcome.  
DELETE | `/api/scores/{score_id}` | Delete one persisted score run.  
POST | `/api/scores/batch` | Start asynchronous scoring for up to 20 profiles.  
GET | `/api/scores/batch/{batch_id}` | Poll an in-memory batch job.

The server also exposes the static test UI at `/` and FastAPI-generated documentation at `/docs`, `/redoc`, and `/openapi.json`.

### Common Behavior

- A validation failure returns 422 Unprocessable Entity using FastAPI's standard detail error array. This applies to malformed JSON, missing required fields, invalid path/query types, and declared field constraints.
- Endpoint-specific failures return an object shaped as `{ "detail": "message" }`.
- Timestamps use ISO 8601 UTC strings when present.
- Scores are persisted in the local SQLite database configured by `DATABASE_PATH` (default: `data/scores.db`).
- `red_flag.status` is one of `pass`, `review`, or `fail`. It is a separate screening gate and is not included in the composite-score calculation.
- Human-review statuses are `pending`, `approved`, `rejected`, and `overridden`.
- Candidate imports are stored in the separate candidates table. Importing a profile does not create a score run or queue scoring.

## Shared Schemas

### CandidateProfileInput

Used by POST /api/score, POST /api/scores/batch, and each profiles item in POST /api/candidates/import.

Field | Type | Required | Constraints and meaning

candidate\_label | string | Yes | Trimmed non-blank label; 1-200 characters. Use a non-identifying label such as sample\_001.

job\_role | string | No | Defaults to ""; maximum 200 characters. Target role used as scoring context.

cv\_claims | string | No | Defaults to ""; maximum MAX\_FIELD\_CHARS characters (8,000 by default). Resume/CV claims used for consistency checks.

profile\_about | string | No | Defaults to ""; maximum MAX\_FIELD\_CHARS characters. Public bio/about text.

posts\_sample | string | No | Defaults to ""; maximum MAX\_FIELD\_CHARS characters. Plain-text sample of posts or articles.

comments\_sample | string | No | Defaults to ""; maximum MAX\_FIELD\_CHARS characters. Plain-text sample of comments or interactions.

network\_notes | string | No | Defaults to ""; maximum MAX\_FIELD\_CHARS characters. Notes about endorsements or professional connections.

MAX\_FIELD\_CHARS can be overridden with an environment variable. Text fields are trimmed before scoring.

Example:

```
{
  "candidate_label": "sample_001",
  "job_role": "Backend Engineer",
  "cv_claims": "Software Engineer, 2022-2024. Python and SQL.",
  "profile_about": "Backend engineer interested in distributed systems.",
  "posts_sample": "Wrote about database indexing strategies.",
  "comments_sample": "Helpful technical comments on engineering posts.",
  "network_notes": "Several relevant endorsements."
}
```

### WebhookRegisterRequest and WebhookResponse

Used by POST /api/webhooks.

Request field | Type | Required | Constraints and meaning

url | string | Yes | Trimmed HTTP or HTTPS URL with a host; 1-2,048 characters.

events | string array | No | Defaults to []; maximum 20 items. Values are trimmed and cannot be blank.

WebhookResponse returns the persisted id (integer), url (string), events (string array), and created\_at (ISO 8601 UTC string).

### TokenRequest and TokenResponse

Used by POST /api/auth/token.

Request field | Type | Required | Constraints and behavior

client\_id | string | Yes | Required identifier; 1-200 characters. The current prototype validates its presence and length but does not look it up.

client\_secret | string | Yes | Required shared secret; 1-512 characters. It must match the configured API\_SECRET\_KEY.

TokenResponse contains access\_token (string) and token\_type (always "bearer"). The current prototype returns the configured shared secret itself

as access\_token; it does not issue unique tokens, expiry times, or refresh tokens. Treat this endpoint as local development-only functionality and do not log either credential or token.

## CandidateImportRequest and CandidateImportResponse

Used by POST /api/candidates/import. The bridge posts its normalized scraper output in the profiles field.

Request field | Type | Required | Constraints and behavior

profiles | array of CandidateProfileInput | Conditionally | Provide a non-empty list of up to 100 profiles. Each item follows CandidateProfileInput validation. Do not supply this field with json\_file.

json\_file | string | Conditionally | Provide JSON text containing either a non-empty profile array or { "profiles": [...] }. Maximum length is MAX\_FIELD\_CHARS \* 100 (800,000 by default). Do not supply this field with profiles.

CandidateImportResponse contains imported\_count (integer), candidate\_ids (integer array), and status (always "success"). IDs identify records in the local candidates table, not score-run IDs.

## ScoreResult

Returned by successful scoring, score detail, review update, comparison, and completed batch items.

Field | Type | Description

id | integer or null | SQLite score-run identifier. It is populated after persistence.

candidate\_label | string | Label supplied in the request.

job\_role | string | Target role context; empty when not supplied.

rubric\_version | string | Human-managed rubric version used for the run.

rubric\_hash | string | Twelve-character fingerprint of rubric content at scoring time.

model\_used | string | Configured Ollama model name.

overall\_summary | string | Model-generated plain-language profile summary.

dimension\_scores | array of DimensionScore | One result for each active rubric dimension.

composite\_score | number | Weighted 0-100 score, rounded to one decimal.

red\_flag | RedFlagResult | Separate pass/review/fail screening result.

human\_review | HumanReviewInfo | Human-review state; initially pending.

excluded\_attributes\_detected | string array | Protected attributes detected but explicitly excluded from scoring.

raw\_model\_output | string | Unprocessed JSON text returned by the local model.

created\_at | string or null | UTC persistence timestamp.

DimensionScore contains key (string), label (string), score (number),

and rationale (string). RedFlagResult contains status and rationale.

HumanReviewInfo contains status, notes, and reviewed\_at (an ISO 8601 string or null).

Example abbreviated response:

```
{
  "id": 42,
  "candidate_label": "sample_001",
  "job_role": "Backend Engineer",
  "rubric_version": "0.2.0",
  "rubric_hash": "abc123def456",
  "model_used": "llama3.1",
  "overall_summary": "Consistent technical profile with clear communication.",
  "dimension_scores": [
    {
      "key": "professional_consistency",
      "label": "Professional Consistency",
      "score": 82.0,
```

```

    "rationale": "Claims align with supplied profile text."
  }
],
"composite_score": 80.5,
"red_flag": { "status": "pass", "rationale": "No concerns found." },
"human_review": { "status": "pending", "notes": "", "reviewed_at": null },
"excluded_attributes_detected": [],
"raw_model_output": "{...}",
"created_at": "2026-07-27T12:34:56.789012+00:00"
}

```

## Endpoints

### GET /api/health

Reports application health and whether Ollama is reachable. The endpoint returns 200 OK even when Ollama is unavailable; inspect `ollama.reachable`.

Parameters: none.

Success response (200 OK):

```

{
  "backend": "ok",
  "ollama": {
    "reachable": true,
    "configured_model": "llama3.1",
    "model_available_locally": true,
    "installed_models": ["llama3.1:latest"]
  }
}

```

When the Ollama check fails, `ollama` instead contains `reachable: false`, `configured_model`, and an error string.

### GET /api/rubric

Returns the rubric currently used by scoring. It lets clients display the dimensions, weights, exclusion policy, and version that govern new runs.

Parameters: none.

Success response (200 OK):

| Field                            | Type         | Description                                                                                       |
|----------------------------------|--------------|---------------------------------------------------------------------------------------------------|
| <code>dimensions</code>          | array        | Objects with key, label, weight, and description. Current weights are 0.33, 0.27, 0.27, and 0.13. |
| <code>red_flag_note</code>       | string       | Explains that red-flag screening is a separate escalation gate.                                   |
| <code>excluded_attributes</code> | string array | Attributes that must not influence scores.                                                        |
| <code>rubric_version</code>      | string       | Current configured version, 0.2.0.                                                                |
| <code>rubric_hash</code>         | string       | Current twelve-character rubric-content fingerprint.                                              |

### POST /api/score

Scores one profile through the configured local Ollama model and persists the result in SQLite.

Request body: a `CandidateProfileInput` <#candidateprofileinput> object.

Success response (200 OK): a persisted `ScoreResult` <#scorerresult>, including `id` and `created_at`.

Error responses:

| Status | When returned                                                        |
|--------|----------------------------------------------------------------------|
| 422    | The input body does not satisfy <code>CandidateProfileInput</code> . |

502 | The scoring pipeline cannot call or parse the configured Ollama model. The response detail contains the scoring error.

## **POST /api/webhooks**

Registers a webhook destination and the event names it is interested in. The route persists the registration locally; event delivery is not implemented by the current prototype.

Request body: a

WebhookRegisterRequest <#webhookregisterrequest-and-webhookresponse> object.

Example:

```
{
  "url": "https://example.test/hooks/behavior-score",
  "events": ["score.completed", "score.review_required"]
}
```

Success response (201 Created):

```
{
  "id": 7,
  "url": "https://example.test/hooks/behavior-score",
  "events": ["score.completed", "score.review_required"],
  "created_at": "2026-07-27T12:34:56.789012+00:00"
}
```

Error response: 422 Unprocessable Entity when url is blank, is not an HTTP(S) URL with a host, events has more than 20 values, or an event value is blank.

## **POST /api/auth/token**

Exchanges client credentials for a development bearer token. The server compares client\_secret with API\_SECRET\_KEY using a constant-time comparison. The required client\_id is currently not used for identity lookup.

Request body: an

TokenRequest <#tokenrequest-and-tokenresponse> object.

Example:

```
{
  "client_id": "local-scraper-bridge",
  "client_secret": "configured-api-secret"
}
```

Success response (200 OK): an

TokenResponse <#tokenrequest-and-tokenresponse> object.

Error responses:

Status | When returned

401 | client\_secret does not match the configured shared secret. The response includes WWW-Authenticate: Bearer.

422 | The body is malformed, a credential is missing, or client\_id/client\_secret violates its length constraint.

The returned token is not currently required by any route, has no expiry, and is not an authentication boundary. Do not expose this prototype endpoint beyond a trusted local development environment.

## **POST /api/candidates/import**

Persists normalized profiles from an approved scraper bridge into the local candidates table. It does not score the profiles or create a batch job. Submit only sample, synthetic, or explicitly consented data with a valid legal basis; this route must not be used to ingest scraped candidate data indiscriminately.

Request body: a CandidateImportRequest <#candidateimportrequest> object.

Example:

```
{
  "profiles": [
    {
      "candidate_label": "consented_sample_001",
      "job_role": "Backend Engineer",
      "cv_claims": "Software Engineer, 2022-2024. Python and SQL.",
      "profile_about": "Backend engineer interested in distributed systems.",
      "posts_sample": "Wrote about database indexing strategies.",
      "comments_sample": "Helpful technical comments on engineering posts.",
      "network_notes": "Several relevant endorsements."
    }
  ]
}
```

Success response (200 OK):

```
{
  "imported_count": 1,
  "candidate_ids": [17],
  "status": "success"
}
```

candidate\_ids contains the newly persisted local candidate-record IDs. They cannot be passed to score-detail or batch-status routes, and no scoring is started by this endpoint.

The same import can be submitted from JSON-file contents:

```
{
  "json_file": "{\"profiles\": [{\"candidate_label\": \"consented_sample_001\"}]}"
}
```

Error responses:

Status | When returned

422 | Neither or both import sources are provided, the supplied list is empty or has more than 100 entries, an individual profile fails CandidateProfileInput validation, or json\_file is invalid or exceeds its size limit.

scripts/scraper\_bridge.py accepts a JSON array, a single candidate object, or an envelope using candidates or profiles. It normalizes common scraper field names into the profiles request contract and submits no more than 100 profiles per import request.

## GET /api/scores

Lists persisted score runs. Filters are combined with logical AND. The search filter matches candidate\_label, job\_role, and overall\_summary.

Query parameter | Type | Default | Constraints and behavior

limit | integer | 50 | Minimum 1; maximum 200.

offset | integer | 0 | Minimum 0.

search | string | null | Case handling is delegated to SQLite LIKE; surrounding whitespace is ignored.

red\_flag\_status | string | null | Stored red-flag status to match, normally pass, review, or fail.

human\_review\_status | string | null | Stored review status to match: pending, approved, rejected, or overridden.

min\_score | number | null | Inclusive minimum composite score; 0-100.

max\_score | number | null | Inclusive maximum composite score; 0-100.

has\_excluded\_attributes | boolean | null | true returns runs with detected excluded attributes; false returns runs without any.

sort\_by | string | id | One of id, composite\_score, created\_at, or candidate\_label.

sort\_order | string | desc | asc or desc.

Success response (200 OK):

```
{
  "results": ["ScoreResult objects"],
  "total": 42,
  "limit": 50,
  "offset": 0,
  "filters": {
    "search": null,
    "red_flag_status": null,
    "human_review_status": null,
    "min_score": null,
    "max_score": null,
    "has_excluded_attributes": null,
    "sort_by": "id",
    "sort_order": "desc"
  }
}
```

total is the count before pagination but after all filters. The results items use the same shape as ScoreResult <#scoreresult>.

## GET /api/scores/analytics

Returns aggregate metrics for every persisted score run. It accepts no query parameters.

Success response (200 OK):

Field | Type | Description

total\_candidates | integer | Total number of persisted runs.

composite\_score\_stats | object | avg, median, min, and max numeric composite scores. All are 0.0 when no runs exist.

red\_flag\_breakdown | object | Counts keyed by pass, review, and fail.

human\_review\_breakdown | object | Counts keyed by review status.

excluded\_attributes\_counts | object | Count of each detected excluded-attribute label.

score\_buckets | object | Counts in 0\_to\_20, 21\_to\_40, 41\_to\_60, 61\_to\_80, and 81\_to\_100.

## GET /api/scores/export

Downloads up to 10,000 filtered score runs. Export results are always ordered by score ID descending. The has\_excluded\_attributes list filter is not available on this route.

Query parameter | Type | Default | Constraints and behavior

format | string | csv | csv or json.

search | string | null | Same matching behavior as GET /api/scores.

red\_flag\_status | string | null | Same filtering behavior as GET /api/scores.

human\_review\_status | string | null | Same filtering behavior as GET /api/scores.

min\_score | number | null | Inclusive composite-score lower bound.

max\_score | number | null | Inclusive composite-score upper bound.

Success responses (200 OK):

format | Content type | Download name | Body

csv | text/csv | candidate\_scores\_export.csv | Header row followed by one row per score.  
Includes review fields and flattened red-flag data.  
json | application/json | candidate\_scores\_export.json | JSON array of ScoreResult  
<#scoreresult>-shaped score runs.

The download filename is supplied through the Content-Disposition response header.

## GET /api/scores/compare

Compares the selected persisted score runs. Candidate results follow the order of the requested IDs when those IDs exist.

Query parameter | Type | Required | Description  
ids | string | Yes | Comma-separated integer score IDs, for example 1,2,5.

Success response (200 OK):

Field | Type | Description  
candidates | array of ScoreResult | Score runs found for the supplied IDs.  
dimension\_averages | object | Average score by dimension key, rounded to one decimal.  
highest\_scoring\_candidate | string or null | Candidate label with the highest composite score.  
lowest\_scoring\_candidate | string or null | Candidate label with the lowest composite score.  
red\_flags\_summary | object | Counts keyed by pass, review, and fail.

Error responses:

Status | When returned  
400 | ids is not a comma-separated integer list or contains no IDs.  
404 | None of the requested IDs exist.

## GET /api/scores/{score\_id}

Returns one persisted scoring run in full detail.

Path parameter | Type | Description  
score\_id | integer | SQLite score-run ID.

Success response (200 OK): a ScoreResult <#scoreresult> object.

Error response: 404 Not Found with { "detail": "Score not found" } when the ID is absent.

## PATCH /api/scores/{score\_id}/review

Records a human-review outcome and notes for an existing score run. The server sets human\_review.reviewed\_at to the update time.

Path parameter | Type | Description  
score\_id | integer | SQLite score-run ID.

Request body:

Field | Type | Required | Constraints  
status | string | Yes | One of pending, approved, rejected, or overridden.  
notes | string | No | Defaults to ""; maximum 1,000 characters.

Example:

```
{
  "status": "approved",
  "notes": "Reviewer confirmed the evidence is appropriate for this prototype run."
}
```

Success response (200 OK): the updated ScoreResult <#scoreresult>.

Error responses:

Status | When returned

404 | The score-run ID does not exist.

422 | The body has an invalid review status or notes exceed 1,000 characters.

### **DELETE /api/scores/{score\_id}**

Deletes one persisted score run permanently.

Path parameter | Type | Description

score\_id | integer | SQLite score-run ID.

Success response (200 OK):

```
{
  "message": "Score run deleted",
  "id": 42
}
```

Error response: 404 Not Found with

{ "detail": "Score run not found" } when the ID is absent.

### **POST /api/scores/batch**

Creates an in-memory batch job and queues profile scoring as a FastAPI background task. Each successful item is persisted as an individual score run.

Request body:

Field | Type | Required | Constraints

profiles | array of CandidateProfileInput | Yes | At least 1 and at most 20 profile objects.

Example:

```
{
  "profiles": [
    { "candidate_label": "sample_001", "cv_claims": "Developer" },
    { "candidate_label": "sample_002", "cv_claims": "Analyst" }
  ]
}
```

Success response (200 OK): a BatchJobStatus <#batchjobstatus> object. A new response starts with status: "processing" and empty results; poll the batch-status endpoint for completion.

Errors: 422 Unprocessable Entity when profiles is outside the allowed size or an item does not satisfy CandidateProfileInput.

### **GET /api/scores/batch/{batch\_id}**

Returns the latest state of an in-memory batch job.

Path parameter | Type | Description

batch\_id | string | Eight-character ID returned when the batch was created.

### **BatchJobStatus**

Field | Type | Description

batch\_id | string | Batch identifier.

status | string | processing, completed, or failed.

total\_items | integer | Number of profiles accepted in the request.

completed\_items | integer | Number of successfully scored profiles.

failed\_items | integer | Number of profiles that failed scoring.

results | array | Completed/failed item records accumulated as processing proceeds.

Each item record has candidate\_label, status (completed or failed), score\_result (a ScoreResult or null), and error (a string or null).  
A batch is marked completed if at least one item succeeds; it is marked failed only when every item fails.

Success response (200 OK): a BatchJobStatus <#batchjobstatus> object.

Error response: 404 Not Found with  
{ "detail": "Batch job not found" } when the ID is unknown. Batch state is not durable and is lost when the application process restarts.

## Non-API Routes

Path | Response | Purpose

/ | Static HTML application | Minimal browser test interface served from static/.

/docs | Swagger UI | Interactive FastAPI documentation.

/redoc | ReDoc | Alternative generated documentation interface.

/openapi.json | OpenAPI JSON | Machine-readable API schema.

## Operational Notes

- Configure the local model with OLLAMA\_MODEL; the placeholder value REPLACE\_WITH\_YOUR\_MODEL intentionally prevents scoring.
- GET /api/health is the recommended readiness check before calling score endpoints.
- The API is configured for local prototype use. It does not implement authentication, rate limiting, durable background-job storage, or production-grade error handling.

Source: docs/project\_report.md

# Behavior Scoring ? Comprehensive Project Report

- > Project: Candidate Social Media Behavior Scoring System
- > Version: 0.2.0
- > Branch: API
- > Author: Adrie (Internship ? Semester Pendek)
- > Date: July 27, 2026
- > Stack: Python 3.11 • FastAPI • Ollama (Local LLM) • SQLite • Playwright

## 1. Executive Summary

This repository contains an internal R&D candidate screening & behavioral evaluation system. Candidate social media profile text (from consented, synthetic, or scraped sources) is evaluated against a deterministic, weighted rubric using a local LLM via Ollama.

### Core Pillars

- Privacy-First: 100% local processing via Ollama ? zero candidate data leaves your machine.
- Fairness Protection: 7 excluded protected attributes (religion, ethnicity, politics, marital status, health, sexual orientation, age) are ignored during scoring via prompt engineering AND a deterministic Python keyword backstop.
- Auditable: Every scoring run records the exact model used, rubric version, rubric content hash, raw LLM output, and timestamp.
- Human-in-the-Loop: All AI scores default to human\_review\_status = "pending" requiring recruiter review.

## 2. System Architecture

```
graph TB
  subgraph Frontend
    UI["Static Test UI<br/>(http://127.0.0.1:8000)"]
  end

  subgraph Backend["FastAPI Backend (API Branch)"]
    MAIN["app/main.py<br/>15+ REST API Endpoints"]
    SCORING["app/scoring.py<br/>Prompt Builder + JSON Parser"]
    RUBRIC["app/rubric.py<br/>Rubric Dimensions & Weights"]
    AUTH["app/auth.py<br/>Authentication Middleware"]
    DB["app/db.py<br/>SQLite Storage Layer"]
    SCHEMAS["app/schemas.py<br/>Pydantic Request/Response Models"]
  end

  subgraph Integration
    BRIDGE["scripts/scrapper_bridge.py<br/>Playwright Scrapper Bridge"]
    E2E["tests/e2e/<br/>Playwright E2E Test Suite"]
  end

  subgraph External
    OLLAMA["Ollama LLM Server<br/>localhost:11434"]
    SQLITE["SQLite Database<br/>data/scores.db"]
  end

  UI -->|HTTP| MAIN
  BRIDGE -->|POST /api/candidates/import| MAIN
  MAIN --> SCORING
  SCORING --> RUBRIC
  SCORING -->|httpx async| OLLAMA
  MAIN --> DB
  DB --> SQLITE
  MAIN --> SCHEMAS
  MAIN --> AUTH
  E2E -->|Playwright APIRequestContext| MAIN
```

## 3. Implemented API Contract

| Category  | Endpoint                     | Method | Description                                                                   |
|-----------|------------------------------|--------|-------------------------------------------------------------------------------|
| System    | /api/health                  | GET    | Reports backend, Ollama server, and local LLM model reachability.             |
|           | /api/rubric                  | GET    | Returns rubric dimensions, weights, red flag rules, and content hash.         |
| Auth      | /api/auth/token              | POST   | Issues Bearer access token for valid client credentials.                      |
| Scoring   | /api/score                   | POST   | Scores a candidate profile via Ollama and saves result to SQLite.             |
|           | /api/scores                  | GET    | Returns paginated scores with text search, score range, and red-flag filters. |
|           | /api/scores/{id}             | GET    | Fetches complete scoring breakdown and raw LLM output by ID.                  |
|           | /api/scores/{id}             | DELETE | Removes a candidate score entry.                                              |
| Batch     | /api/scores/batch            | POST   | Enqueues a batch of profiles for background scoring.                          |
|           | /api/scores/batch/{batch_id} | GET    | Checks processing status and item results of a batch job.                     |
| Import    | /api/candidates/import       | POST   | Bulk-imports scraped candidate profiles into SQLite.                          |
| Webhooks  | /api/webhooks                | POST   | Registers webhook endpoint URLs for completion event notifications.           |
| Review    | /api/scores/{id}/review      | PATCH  | Updates recruiter audit status (approved, rejected, overridden) and notes.    |
| Analytics | /api/scores/analytics        | GET    | Returns aggregated composite score stats, red flags, and score buckets.       |
|           | /api/scores/export           | GET    | Exports scored candidate data as CSV or JSON download.                        |
|           | /api/scores/compare          | GET    | Side-by-side comparison of candidate profiles with dimension averages.        |

## 4. Test Verification Summary

The test suite contains 88 automated tests spanning unit, contract integration, and Playwright E2E tests:

```
# Run the test suite:  
venv\Scripts\python.exe -m pytest
```

### Test Breakdown

- Playwright E2E Suite (tests/e2e/): 18 tests (APIRequestContext testing health, scoring, auth token issuance, candidate import, webhooks).
- Backend Feature & Contract Suite (tests/integration/): 18 tests (SQLite persistence, analytics, filters, pagination, export).
- Rubric & Scoring Pipeline Unit Suite (tests/unit/): 52 tests (dimension weight validation, non-numeric fallback handling, excluded-attribute keyword backstop).

Result: 88 passed (100% pass rate).

## 5. Single Workspace Source Control Workflow

All work is maintained directly on the API branch:

- No extra Git worktrees or temporary branches cluttering Source Control.
- Direct commit and push to origin/API.